

1 Regression

1.1 Linear Regression

Simple $y_n \approx f(\mathbf{x}_n) := w_0 + w_1 x_{n1}$

Multiple

$$f(\mathbf{x}_n) := w_0 + \sum_{j=1}^D w_j x_{nj} = \tilde{\mathbf{x}}_n^T \mathbf{w}$$

If $D > N$ the task is underdetermined (more dimensions than data) \rightarrow regularization.

2 Cost functions

MSE = $\frac{1}{N} \sum_{n=1}^N [y_n - f(\mathbf{x}_n)]^2$ Not good with outliers.

$$\text{MAE} = \frac{1}{N} \sum_{n=1}^N |y_n - f(\mathbf{x}_n)|$$

Error $e_n = y_n - f(\mathbf{x}_n)$

2.1 Convexity

A line joining two points never intersects with the function anywhere else. $f(\lambda \mathbf{u} + (1 - \lambda)\mathbf{v}) \leq \lambda f(\mathbf{u}) + (1 - \lambda)f(\mathbf{v})$ with $\lambda \in [0; 1]$. A strictly convex function has a unique global minimum \mathbf{w}^* . Sums of convex functions are convex.

A function must always lie above its linearisation:

$$\mathcal{L}(u) \geq \mathcal{L}(w) + \nabla \mathcal{L}(w)^T (u - w) \forall u, w.$$

A set is convex iff line segment between any two points of \mathcal{C} lies in \mathcal{C} :

$$\theta u + (1 - \theta)v \in \mathcal{C}$$

3 Optimisation

Find $\mathbf{w}^* \in \mathcal{R}^D$ which $\min \mathcal{L}(\mathbf{w})$.

$$\text{Gradient } \nabla \mathcal{L} := \begin{bmatrix} \frac{\partial \mathcal{L}(\mathbf{w})}{\partial w_1} & \dots & \frac{\partial \mathcal{L}(\mathbf{w})}{\partial w_D} \end{bmatrix}$$

3.1 Gradient descent

$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - \gamma \nabla \mathcal{L}(\mathbf{w}^{(t)})$. Very sensitive to ill-conditioning.

GD - Linear Reg

$$\mathcal{L}(\mathbf{w}) = \frac{1}{2N} (\mathbf{y} - X\mathbf{w})^T (\mathbf{y} - X\mathbf{w}) \rightarrow$$

$$\nabla \mathcal{L}(\mathbf{w}) = -\frac{1}{N} X^T (\mathbf{y} - X\mathbf{w}). \text{ Cost:}$$

$$O_{err} = 2ND + N \text{ and } O_w = 2ND + D.$$

3.2 SGD

$$\mathcal{L} = \frac{1}{N} \sum \mathcal{L}_n(\mathbf{w}) \text{ with update } \mathbf{w}^{(t+1)} =$$

$$\mathbf{w}^{(t)} - \gamma \nabla \mathcal{L}_n(\mathbf{w}^{(t)}).$$

3.3 Mini-batch SGD

$$\mathbf{g} = \frac{1}{|B|} \sum_{n \in B} \nabla \mathcal{L}_n(\mathbf{w}^{(t)}) \text{ with update}$$

$$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - \gamma \mathbf{g}.$$

3.4 Subgradient at w

$\mathbf{g} \in \mathcal{R}^D$ such that $\mathcal{L}(u) \geq \mathcal{L}(w) + \mathbf{g}^T (u - w)$. Example subgradient for MAE: $h(e) = |e| \rightarrow g(e) = \text{sgn}(e)$ if $e \neq 0, \lambda$ otherwise. We get the gradient:

$$\nabla \mathcal{L}_{MAE} = -\frac{1}{N} \sum_n \text{sgn}(x_n) \nabla f(x_n).$$

2.1 Subgradients

$$\mathbf{w}^{(t+1)} = \mathcal{P}_{\mathcal{C}}[\mathbf{w}^{(t)} - \gamma \nabla \mathcal{L}(\mathbf{w}^{(t)})]$$

3.6 Newton's method

Second order (more expensive $O(ND^2 + D^3)$ but faster convergence).

$$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - \gamma^{(t)} (H^{(t)})^{-1} \nabla \mathcal{L}(\mathbf{w}^{(t)})$$

3.7 Optimality conditions

Necessary : $\nabla \mathcal{L}(\mathbf{w}^*) = 0$ Sufficient :

$$\text{Hessian PSD } \mathbf{H}(\mathbf{w}^*) := \frac{\partial^2 \mathcal{L}(\mathbf{w}^*)}{\partial w \partial w^T}$$

4 Least Squares

4.1 Normal Equation

$$X^T (\mathbf{y} - X\mathbf{w}) = 0 \Rightarrow$$

$\mathbf{w}^* = (XX^T)^{-1} X^T \mathbf{y}$ and $\hat{\mathbf{y}}_m = \mathbf{x}_m^T \mathbf{w}^*$
Graham matrix invertible iff $\text{rank}(X) = D$ (use SVD $X = USV^T$ if this is not the case to get pseudo-inverse $\mathbf{w}^* = V\hat{S}U^T$ with \hat{S} pseudo-inverse of S).

5 Likelihood

Probabilistic model $y_n = \mathbf{x}_n^T \mathbf{w} + \epsilon_n$. Probability of observing the data given a set of parameters and inputs : $p(\mathbf{y}|X, \mathbf{w}) = \prod p(y_n | \mathbf{x}_n, \mathbf{w}) = \prod \mathcal{N}(y_n | \mathbf{x}_n^T \mathbf{w}, \sigma^2)$

Best model maximises log-likelihood

$$\mathcal{L}_{LL} = -\frac{1}{2\sigma^2} \sum (y_n - \mathbf{x}_n^T \mathbf{w})^2 + cst.$$

6 Regularization

6.1 Ridge Regression

$$\mathcal{L}(\mathbf{w}) = \frac{1}{2} (\mathbf{y} - X\mathbf{w})^T (\mathbf{y} - X\mathbf{w}) + \frac{\lambda}{2} \|\mathbf{w}\|_2^2 \rightarrow$$

$$\mathbf{w}_{\text{ridge}}^* = (XX^T + \lambda I_D)^{-1} X^T \mathbf{y} = X^T (XX^T + \lambda I_N)^{-1} \mathbf{y}$$

Can be considered a MAP estimator : $\mathbf{w}_{\text{ridge}}^* = \text{argmin}_w - \log(p(w|X, y))$

6.2 Lasso

$$\text{Sparse solution. } \mathcal{L}(w) = \frac{1}{2N} (y - Xw)^T (y - Xw) + \lambda \|w\|_1$$

7 Model Selection

7.1 Bias-Variance decomposition

Small dimensions : large bias, small variance. Large dimensions : small bias, large variance. Error for the val set compared to the emp distr of the data goes down like $\frac{1}{\sqrt{|\text{validation points}|}}$ and

goes up like $\sqrt{\ln(|\text{hyper parameters}|)}$

8 Classification

8.1 Optimal

$$\hat{y}(\mathbf{x}) = \text{argmax}_{v \in \mathcal{Y}} p(y|\mathbf{x})$$

8.2 Logistic regression

$\sigma(z) = \frac{e^z}{1+e^z}$ to limit the predicted values $y \in [0; 1]$ ($p(1|\mathbf{x}) = \sigma(\mathbf{x}^T \mathbf{w})$ and $p(0|\mathbf{x}) = 1 - \sigma(\mathbf{x}^T \mathbf{w})$). We decide with respect to 0.5

Likelihood

$$p(y|X, w) = \prod_{n: y_n=0} p(y_n = 0 | x_n) \dots \prod_{n: y_n=K} p(y_n = K | x_n) = \prod_k^K \prod_n [p(y_n = k | x_n, w)]^{\tilde{y}_{nk}}$$

where $\tilde{y}_{nk} = 1$ if $y_n = k$.

For binary classification

$$p(y|X, w) = \prod_{n: y_n=0} p(y_n = 0 | x_n) \prod_{n: y_n=1} p(y_n = 1 | x_n) = \prod_n \sigma(x_n^T w)^{y_n} [1 - \sigma(x_n^T w)]^{1-y_n}$$

Loss

$$\mathcal{L}(w) = \sum_{n=1}^N \ln(1 + \exp(x_n^T w)) - y_n x_n^T w$$

which is convex in w .

Gradient

$$\nabla \mathcal{L}(w) = \sum_{n=1}^N x_n (\sigma(x_n^T w) - y_n) = X^T [\sigma(Xw) - y]$$

(no closed form solution).

Hessian

$$H(w) = X^T S X \text{ with } S_{nn} = \sigma(x_n^T w) [1 - \sigma(x_n^T w)]$$

8.3 Exponential family

General form

$$p(y|\eta) = h(y) \exp[\eta^T \psi(y) - A(\eta)] \text{ where}$$

Cumulant

$$A(\eta) = \ln \left[\int_{\mathcal{Y}} h(y) \exp[\eta^T \psi(y)] dy \right]$$

$$\nabla A(\eta) = \mathbb{E}[\psi(y)]$$

$$\nabla^2 A(\eta) = \mathbb{E}[\psi \psi^T] - \mathbb{E}[\psi] \mathbb{E}[\psi^T]$$

Link function

$$\eta = g^{-1}(\mu) \Leftrightarrow \mu = g(\eta)$$

$$\bullet \eta_{\text{gaussian}} = (\mu / \sigma^2, -1/2\sigma^2)$$

$$\bullet \eta_{\text{poisson}} = \ln(\mu)$$

$$\bullet \eta_{\text{bernoulli}} = \ln(\mu / (1 - \mu))$$

$$\bullet \eta_{\text{general}} = g^{-1}(\frac{1}{N} \sum_{n=1}^N \psi(y_n))$$

8.4 Nearest Neighbor Models

Performs best in low dimensions.

8.4.1 k-NN

$f_{St,k}(x) = \frac{1}{k} \sum_{n: x_n \in \text{ngb}_{St,k}(x)} y_n$ Pick odd k so there is a clear winner. Large $k \rightarrow$ large bias small variance (inv.)

8.4.2 Error bound

$$\mathbb{E}[\mathcal{L}_{St}] \leq 2\mathcal{L}_{f^*} + 4c\sqrt{d}N^{-1/d+1}$$

8.5 Support Vector Machines (SVM)

Logistic regression with hinge loss : $\min_w \sum_{n=1}^N [1 - y_n x_n^T w]_+ + \frac{\lambda}{2} \|w\|^2$ where $y \in [-1; 1]$ is the label and $\text{hinge}(x) = \max\{0, x\}$. Convex but not differentiable so need subgradient. We can also use duality : $\mathcal{L}(w) = \max_{\alpha} G(w, \alpha)$. For SVM $\min_w \max_{\alpha \in [0, 1]^N} \sum \alpha_n (1 - y_n x_n^T w) + \frac{\lambda}{2} \|w\|^2$ differentiable and convex. Can switch \max and \min when convex in w and concave in α . This can make the formulation simpler:

$w(\alpha) = \frac{1}{\lambda} \sum \alpha_n y_n x_n = \frac{1}{\lambda} X^T \text{diag}(y) \alpha$ which yields the optimisation problem: $\max_{\alpha \in [0, 1]^N} \alpha^T \mathbf{1} - \frac{1}{2\lambda} \alpha^T Y X X^T Y \alpha$ The solution is sparse (α_n is the slope of the lines that are lower bounds to the hinge loss).

8.6 Kernel Ridge Regression

From duality $w^* := X^T \alpha^*$ where $\alpha^* := (K + \lambda I_N)^{-1} y$ and $K = XX^T = \phi^T(x) \phi(x) = \kappa(x, x')$ (needs to be PSD and symmetric).

9 Unsupervised Learning

9.1 K-means clustering

$$\min \mathcal{L}(z, \mu) = \sum_n \sum_k^K z_{nk} \|x_n - \mu_k\|_2^2 \text{ with } z_{nk} \in \{0, 1\} \text{ (unique assignments: } \sum_k z_{nk} = 1).$$

Algorithm (Coordinate Descent)

$$1. \forall n \text{ compute } z_n = \begin{cases} 1 & \text{if } k = \text{argmin}_j \|x_n - \mu_j\|^2 \\ 0 & \text{otherwise} \end{cases}$$

$$2. \forall k \text{ compute } \mu_k = \frac{\sum_n z_{nk} x_n}{\sum_n z_{nk}}$$

Issues

- Heavy computation
- Spherical clusters
- Hard clusters

Probabilistic model $p(X|\mu, z) = \prod_n \mathcal{N}(x_n | \mu_k, I) = \prod_n \prod_k^K \mathcal{N}(x_n | \mu_k, I)^{z_{nk}}$

9.2 Gaussian Mixture Models

$$p(X|\mu, z) = \prod_n p(x_n | z_n, \mu_k, \Sigma_k) p(z_n | \pi) = \prod_n \prod_k^K [\mathcal{N}(x_n | \mu_k, \Sigma_k)]^{z_{nk}} \prod_k^K [\pi_k]^{z_{nk}}$$

where $p_i k = p(z_n = k)$

Marginal likelihood: z_n are latent variables so they can be factored

out from the likelihood $p(x_n | \theta) = \sum \pi_k \mathcal{N}(x_n | \mu_k, \Sigma_k)$. (number of parameters reduced from $O(N)$ to $O(D^2 K)$).

9.3 EM

9.3.1 GMM

Intialize $\mu^{(1)}, \Sigma^{(1)}, \pi^{(1)}$.

- E-step: Compute the assignments. $q_{kn}^{(t)} := \frac{\pi_k^{(t)} \mathcal{N}(x_n | \mu_k^{(t)}, \Sigma_k^{(t)})}{\sum_k^K \pi_k^{(t)} \mathcal{N}(x_n | \mu_k^{(t)}, \Sigma_k^{(t)})}$

2. Compute Marginal Likelihood

3. M-step: Update

$$\mu^{(t+1)} = \frac{\sum_n q_{kn}^{(t)} x_n}{\sum_n q_{kn}^{(t)}}$$

$$\Sigma^{(t+1)} = \frac{\sum_n q_{kn}^{(t)} (x_n - \mu^{(t+1)}) (x_n - \mu^{(t+1)})^T}{\sum_n q_{kn}^{(t)}}$$

$$\pi^{(t+1)} = \frac{1}{N} \sum_n q_{kn}^{(t)}$$

9.3.2 General

$$\theta^{(t+1)} := \text{argmax}_{\theta} \sum_n^N \mathbb{E}_{p(z_n | x_n, \theta^{(t)})} [\log p(x_n | \theta)]$$

10 Matrix Factorizations

10.1 Prediction

Find $\mathbf{X} \approx \mathbf{WZ}^T$ where $\mathbf{W} \in \mathbb{R}^{D \times K}$ and $\mathbf{Z} \in \mathbb{R}^{N \times K}$ with $K \ll D, N$. Large $K \rightarrow$ overfitting. If $K \geq \max\{D, N\}$ trivial solution ($\mathbf{W} = \mathbf{1}_D$ or $\mathbf{Z} = \mathbf{1}_N$). Quality of reconstruction (not jointly convex nor identifiable):

$$\mathcal{L}(\mathbf{W}, \mathbf{Z}) := \frac{1}{2} \sum_{(d,n) \in \Omega} [x_{dn} - (\mathbf{WZ}^T)_{dn}]^2 = \sum_{(d,n) \in \Omega} f_{dn}(w, z)$$

$$\text{Regularizer: } \Omega(W, Z) = \frac{\lambda_w}{2} \|\mathbf{W}\|_{Frob}^2 + \frac{\lambda_z}{2} \|\mathbf{Z}\|_{Frob}^2$$

Optimisation with SGD (compute ∇_w for a fixed user d' and ∇_z for a fixed item n'). ALS (assume no missing ratings): $\mathbf{Z}_*^T = (\mathbf{W}^T \mathbf{W} + \lambda_z I_K)^{-1} \mathbf{W}^T \mathbf{X}$ $\mathbf{W}_*^T = (\mathbf{Z}^T \mathbf{Z} + \lambda_w I_K)^{-1} \mathbf{Z}^T \mathbf{X}$

10.2 Text Representation

Factorize the co-occurence matrix to get each row forming a representation of a word (\mathbf{W}) or a context word (\mathbf{Z}) respectively.

10.2.1 GloVe

$$f_{dn} := \min\{1, (n_{dn}/n_{max})^\alpha\}, \alpha \in [0;1]$$

10.2.2 Skipgram/CBOW

Binary classification to separate real word pairs from fake ones.

10.3 FastText

Supervised sentence-level BoW.

11 Dimensionality reduction

11.1 SVD

$\mathbf{X} = \mathbf{USV}^T$, with $\mathbf{X} : D \times N$, $\mathbf{U} : D \times D$ orthonormal, $\mathbf{V} : N \times N$ orthonormal, $\mathbf{S} : D \times N$ diagonal PSD, values in descending order ($s_1 \geq s_2 \geq \dots \geq s_D \geq 0$).

Reconstruction
 $\|\mathbf{X} - \hat{\mathbf{X}}\|_F^2 \geq \|\mathbf{X} - \mathbf{U}_K \mathbf{U}_K^T \mathbf{X}\|_F^2 = \sum_{i \geq K+1} s_i^2 \vee$

rank- K matrix $\hat{\mathbf{X}}$ (i.e. *we should compress the data by projecting it onto these left singular vectors.*)

Truncated SVD: $\mathbf{U}_K \mathbf{U}_K^T \mathbf{X} = \mathbf{US}_K \mathbf{V}^T$

Application to MF: $\mathbf{U} = \mathbf{W}$ and $\mathbf{SV}^T = \mathbf{Z}^T$. Reconstruction limited by the rank- K of \mathbf{W}, \mathbf{Z} .

11.2 PCA

Decorrelate the data. Empirical mean before: $\mathbf{NK} = \mathbf{XX}^T = \mathbf{US}_D^2 \mathbf{U}^T$. After $\tilde{\mathbf{X}} = \mathbf{U}^T \mathbf{X}$: $\mathbf{N}\tilde{\mathbf{K}} = \tilde{\mathbf{X}}\tilde{\mathbf{X}}^T = \mathbf{S}_D^2$ (the components are uncorrelated).

Pitfalls: not invariant under scalings.

12 Quick maff

Chain rule $h = f(g(w)) \rightarrow \partial h(w) = \partial f(g(w)) \nabla g(w)$

Gaussian $\mathcal{N}(y|\mu, \sigma^2) \frac{1}{\sqrt{2\pi\sigma^2}} \exp(-\frac{(y-\mu)^2}{\sigma^2})$

Multivariate Gaussian $\mathcal{N}(y|\mu, \Sigma) \frac{1}{\sqrt{(2\pi)^D \det(\Sigma)}} \exp(-\frac{1}{2}(y - \mu)^T \Sigma^{-1} (y - \mu))$

Bayes rule $p(x|y) = \frac{p(y|x)p(x)}{p(y)}$

Logit $\sigma(x) = \frac{\partial \ln[1+e^x]}{\partial x}$

Naming Joint distribution $p(x, y) = p(x|y)p(y) = p(y|x)p(x)$ where

- $p(x|y) \rightarrow$ likelihood
- $p(y) \rightarrow$ prior
- $p(y|x) \rightarrow$ posterior
- $p(x) \rightarrow$ marginal likelihood

Marginal Likelihood

$$p(\mathbf{X}|\alpha) = \int_{\theta} p(\mathbf{X}|\theta) p(\theta|\alpha) d\theta$$

Posterior probability \propto Likelihood \times Prior

Maximising over a Gaussian is equivalent to minimising MSE:

$$\beta_{MAP}^* = \operatorname{argmax}_{\beta} p(y|\mathbf{X}, \beta) p(\beta) \Leftrightarrow \beta^* = \operatorname{argmin}_{\beta} \mathcal{L}(\beta)$$

Identifiable model $\theta_1 = \theta_2 \rightarrow P_{\theta_1} = P_{\theta_2}$

12.1 Algebra

$$(PQ + I_N)^{-1} P = P(QP + I_M)^{-1}$$

$$\sum_n (y_n - \beta^T \mathbf{x}_n)^2 = (\mathbf{y} - \mathbf{X}\beta)^T (\mathbf{y} - \mathbf{X}\beta)$$

$$\sum_j \beta^2 = \beta^T \beta$$

Unitary / orthogonal: $\mathbf{UU}^T = \mathbf{U}^T \mathbf{U} = \mathbf{I}$ and $\mathbf{U}^T = \mathbf{U}^{-1}$. Rotation matrix (preserves length of vector).

13 Mock Exam Notes

13.1 Normal equation

Unique if convex.

$$\frac{1}{\sigma_k^2} \mathbf{X} (\mathbf{X}^T \mathbf{w}_k - y_k) + \mathbf{w}_k = 0 \Leftrightarrow$$

$$\mathbf{w}_k^* = (\frac{1}{\sigma_k^2} \mathbf{X} \mathbf{X}^T + I_D)^{-1} \frac{1}{\sigma_k^2} \mathbf{X} y_k$$

13.2 MAP solution

$$\mathcal{L}(w) = \sum_k \sum_n \frac{1}{2\sigma_k^2} (y_{nk} - x_n^T w_k)^2 +$$

$$\frac{1}{2} \sum_k \|\mathbf{w}_k\|_2^2 \rightarrow \text{Likelihood } p(y|\mathbf{X}, \mathbf{w}) =$$

$$\prod_n \prod_k \mathcal{N}(y_{nk} | \mathbf{w}_k^T \mathbf{x}_n, \sigma_k^2) \text{ and prior } p(\mathbf{w}) = \prod_k \mathcal{N}(\mathbf{w}_k | 0, I_D)$$

13.3 Convexity

$\ln[\sum_k^K e^{t_k}]$ is convex. Linear sum of parameters is convex.

13.4 Deriving marginal distribution

$$p(y_n | x_n, r_n = k, \beta) = \mathcal{N}(y_n | \beta_k^T \tilde{x}_n, 1)$$

Assume r_n follows a multinomial $p(r_n = k | \pi)$. Derive the marginal $p(y_n | x_n, \beta, \pi)$.
 $p(y_n | x_n, r_n = k, \beta) = \sum_k^K p(y_n, r_n = k | x_n, \beta, \pi) = \sum_k^K p(y_n | r_n = k, x_n, \beta, \pi) \cdot \pi_k = \sum_k^K \mathcal{N}(y_n | \beta_k^T \tilde{x}_n, \sigma^2) \cdot \pi_k$

13.5 MF

$$\hat{r}_{um} = \langle \mathbf{v}_u, \mathbf{w}_m \rangle + b_u + b_m \quad \mathcal{L} = \frac{1}{2} \sum_u m (\hat{r}_{um} - r_{um})^2 + \frac{\lambda}{2} \left[\sum_u (b_u^2 + \|\mathbf{v}_u\|^2) + \sum_m (b_m^2 + \|\mathbf{w}_m\|^2) \right].$$

The optimal value for b_u for a particular user u' : $\sum_{u'} m' (\hat{r}_{u'm} - r_{u'm}) + \lambda b_{u'} = 0$.

Problem jointly convex? Compute $H(\hat{r}(v, w)) = \begin{bmatrix} 2w^2 & 4vw - 2r \\ 4vw - 2r & 2v^2 \end{bmatrix}$

which is not PSD in general.

14 Multiple Choice Notes

14.1 True statements

- Regularisation term sometimes renders the min. problem into a strictly concave/convex problem.

- k-NN can be applied even if the data cannot be linearly separated.

$$\max\{0, x\} = \max_{\alpha \in [0,1]} \alpha x$$

$$\min\{0, x\} = \min_{\alpha \in [0,1]} \alpha x$$

$$g(x) = \min_y f(x, y) \Rightarrow g(x) \leq f(x, y)$$

$$\max_x g(x) \leq \max_x f(x, y)$$

$$\max_x \min_y f(x, y) \leq \min_y \max_x f(x, y)$$

$$\nabla_{\mathbf{W}} (\mathbf{x}^T \mathbf{W} \mathbf{x}) = \mathbf{x} \mathbf{x}^T$$

$$\nabla_{\mathbf{x}} (\mathbf{x}^T \mathbf{W} \mathbf{x}) = (\mathbf{W} + \mathbf{W}^T) \mathbf{x}$$

- If we initialize the K-means algorithm with optimal clusters then it will find in one step optimal representation points.

- If we initialize the K-means algorithm with optimal representation points then it will find in one step optimal clusters.

- Logistic loss is typically preferred over L_2 loss in classification tasks.


- For optimizing a matrix factorization of a $D \times N$ matrix, for large D, N : per iteration, ALS has an increased computational cost over SGD and per iteration, SGD cost is independent of D, N .

- A neural net with one hidden layer and an arbitrary number of hidden nodes with sigmoid activation functions can approximate any “sufficiently smooth” function on a bounded domain.

- The complexity of the back-propagation algorithm for a neural net with L layers and K nodes per layer is $O(K^2 L)$

- Consider a convolutional net where the data is laid out in a one-dimensional fashion and the filter/kernel has M non-zero terms. Ignoring the bias terms, there are M parameters.

14.2 Bayes nets



- X_1 and X_4 are independent.

- X_1 and X_4 are **not** independent given X_6 .

- X_1 and X_4 are independent given X_2 .

- X_1 and X_4 are independent given X_2 and X_3 .

- X_1 and X_4 are independent given X_5 .

14.3 Convex functions

$$f(x) = x^\alpha, x \in \mathbb{R}^+, \forall \alpha \geq 1 \text{ or } \leq 0$$

$$f(x) = -x^3, x \in [-1, 0]$$

$$f(x) = e^{ax}, \forall x, a \in \mathbb{R}$$

$$f(x) = \ln(1/x), x \in \mathbb{R}^+$$

$$f(x) = g(h(x)), x \in \mathbb{R}, g, h \text{ convex and increasing over } \mathbb{R}$$

$$f(x) = ax + b, x \in \mathbb{R}, \forall a, b \in \mathbb{R}$$

$$f(x) = |x|^p, x \in \mathbb{R}, p \geq 1$$

$$f(x) = x \log(x), x \in \mathbb{R}^+$$

14.4 Non-convex functions

$$f(x) = x^3, x \in [-1, 1]$$

$$f(x) = e^{-x/2}, x \in \mathbb{R}$$

$$\sum \mathcal{N}$$

$$\sin(x) \forall x \in \mathbb{R}$$