

1 Regression
1.1 Linear Regression
Simple $y_n \approx f(\mathbf{x}_n) := w_0 + w_1 x_{n1}$
Multiple
 $f(\mathbf{x}_n) := w_0 + \sum_{j=1}^D w_j x_{nj} = \tilde{\mathbf{x}}_n^T \mathbf{w}$
If $D > N$ the task is under-determined (more dimensions than data) \rightarrow regularization.

2 Cost functions
MSE = $\frac{1}{N} \sum_{n=1}^N [y_n - f(\mathbf{x}_n)]^2$ Not good with outliers,
MAE = $\frac{1}{N} \sum_{n=1}^N |y_n - f(\mathbf{x}_n)|$
Error $e_n = y_n - f(\mathbf{x}_n)$

2.1 Convexity
A line joining two points never intersects with the function anywhere else.
 $f(\lambda \mathbf{u} + (1 - \lambda)\mathbf{v}) \leq \lambda f(\mathbf{u}) + (1 - \lambda)f(\mathbf{v})$ with $\lambda \in [0; 1]$. A strictly convex function has a unique global minimum \mathbf{w}^* . Sums of convex functions are convex.

A function must always lie above its linearisation:
 $\mathcal{L}(u) \geq \mathcal{L}(w) + \nabla \mathcal{L}(w)^T (u - w) \forall u, w$.
A set is convex iff line segment between any two points of \mathcal{C} lies in \mathcal{C} :
 $\theta u + (1 - \theta)v \in \mathcal{C}$

3 Optimisation
Find $\mathbf{w}^* \in \mathcal{R}^D$ which $\min \mathcal{L}(\mathbf{w})$.
Gradient $\nabla \mathcal{L} := \begin{bmatrix} \frac{\partial \mathcal{L}(\mathbf{w})}{\partial w_1} & \dots & \frac{\partial \mathcal{L}(\mathbf{w})}{\partial w_D} \end{bmatrix}$

3.1 Gradient descent
 $\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - \gamma \nabla \mathcal{L}(\mathbf{w}^{(t)})$. Very sensitive to ill-conditioning.
GD - Linear Reg
 $\mathcal{L}(\mathbf{w}) = \frac{1}{2N} (\mathbf{y} - X\mathbf{w})^T (\mathbf{y} - X\mathbf{w}) \rightarrow$
 $\nabla \mathcal{L}(\mathbf{w}) = -\frac{1}{N} X^T (\mathbf{y} - X\mathbf{w})$. Cost:
 $O_{err} = 2ND + N$ and $O_w = 2ND + D$.

3.2 SGD
 $\mathcal{L} = \frac{1}{N} \sum \mathcal{L}_n(\mathbf{w})$ with update $\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - \gamma \nabla \mathcal{L}_n(\mathbf{w}^{(t)})$.

3.3 Mini-batch SGD
 $\mathbf{g} = \frac{1}{|B|} \sum_{n \in B} \nabla \mathcal{L}_n(\mathbf{w}^{(t)})$ with update $\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - \gamma \mathbf{g}$.

3.4 Subgradient at w
 $\mathbf{g} \in \mathbb{R}^D$ such that $\mathcal{L}(u) \geq \mathcal{L}(w) + \mathbf{g}^T (u - w)$. Example subgradient for MAE: $h(e) = |e| \rightarrow g(e) = \text{sgn}(e)$ if $e \neq 0, \lambda$ otherwise. We get the gradient:
 $\nabla \mathcal{L}_{MAE} = -\frac{1}{N} \sum_n \text{sgn}(x_n) \nabla f(x_n)$.

3.5 Projected SGD
 $\mathbf{w}^{(t+1)} = \mathcal{P}_{\mathcal{C}}[\mathbf{w}^{(t)} - \gamma \nabla \mathcal{L}(\mathbf{w}^{(t)})]$
3.6 Newton's method
Second order (more expensive $O(ND^2 + D^3)$ but faster convergence).

$w^{(t+1)} = w^{(t)} - \gamma^{(t)} (H^{(t)})^{-1} \nabla \mathcal{L}(w^{(t)})$
3.7 Optimality conditions
Necessary : $\nabla \mathcal{L}(\mathbf{w}^*) = 0$ Sufficient : Hessian PSD $\mathbf{H}(\mathbf{w}^*) := \frac{\partial^2 \mathcal{L}(\mathbf{w}^*)}{\partial w \partial w^T}$

4 Least Squares
4.1 Normal Equation
 $X^T (\mathbf{y} - X\mathbf{w}) = 0 \Rightarrow$
 $\mathbf{w}^* = (XX^T)^{-1} X^T \mathbf{y}$ and $\hat{\mathbf{y}}_m = \mathbf{x}_m^T \mathbf{w}^*$
Graham matrix invertible iff $\text{rank}(X) = D$ (use SVD $X = USV^T$ if this is not the case to get pseudo-inverse $\mathbf{w}^* = V \hat{S} U^T$ with \hat{S} pseudo-inverse of S).

5 Likelihood
Probabilistic model $y_n = \mathbf{x}_n^T \mathbf{w} + \epsilon_n$. Probability of observing the data given a set of parameters and inputs : $p(\mathbf{y}|X, \mathbf{w}) = \prod p(y_n | \mathbf{x}_n, \mathbf{w}) = \prod \mathcal{N}(y_n | \mathbf{x}_n^T \mathbf{w}, \sigma^2)$
Best model maximises log-likelihood
 $\mathcal{L}_{LL} = -\frac{1}{2\sigma^2} \sum (y_n - \mathbf{x}_n^T \mathbf{w})^2 + cst$.

6 Regularization
6.1 Ridge Regression
 $\mathcal{L}(\mathbf{w}) = \frac{1}{2} (\mathbf{y} - X\mathbf{w})^T (\mathbf{y} - X\mathbf{w}) + \frac{\lambda}{2} \|\mathbf{w}\|_2^2 \rightarrow$
 $\mathbf{w}_{\text{ridge}}^* = (XX^T + \lambda I_D)^{-1} X^T \mathbf{y} = X^T (XX^T + \lambda I_N)^{-1} \mathbf{y}$
Can be considered a MAP estimator : $\mathbf{w}_{\text{ridge}}^* = \text{argmin}_w - \log(p(w|X, y))$

6.2 Lasso
Sparse solution. $\mathcal{L}(w) = \frac{1}{2N} (y - Xw)^T (y - Xw) + \lambda \|w\|_1$

7 Model Selection
7.1 Bias-Variance decomposition
Small dimensions : large bias, small variance. Large dimensions : small bias, large variance. Error for the val set compared to the emp distr of the data goes down like $\frac{1}{\sqrt{|\text{validation points}|}}$ and goes up like $\sqrt{\ln(|\text{hyper parameters}|)}$

8 Classification
8.1 Optimal
 $\hat{y}(\mathbf{x}) = \text{argmax}_{v \in \mathcal{Y}} p(y|\mathbf{x})$

8.2 Logistic regression
 $\sigma(z) = \frac{e^z}{1+e^z}$ to limit the predicted values $y \in [0; 1]$ ($p(1|\mathbf{x}) = \sigma(\mathbf{x}^T \mathbf{w})$ and $p(0|\mathbf{x}) = 1 - \sigma(\mathbf{x}^T \mathbf{w})$). We decide with respect to 0.5
Likelihood
 $p(y|X, w) = \prod_{n: y_n=0} p(y_n = 0 | x_n) \dots \prod_{n: y_n=K} p(y_n = K | x_n) = \prod_k^K \prod_n [p(y_n = k | x_n, w)]^{y_{nk}}$ where $\tilde{y}_{nk} = 1$ if $y_n = k$.
For binary classification
 $p(y|X, w) = \prod_{n: y_n=0} p(y_n = 0 | x_n) \prod_{n: y_n=1} p(y_n = 1 | x_n) = \prod_n \sigma(x_n^T w)^{y_n} [1 - \sigma(x_n^T w)]^{1-y_n}$
Loss
 $\mathcal{L}(w) = \sum_{n=1}^N \ln(1 + \exp(x_n^T w)) - y_n x_n^T w$ which is convex in w .
Gradient
 $\nabla \mathcal{L}(w) = \sum_{n=1}^N x_n (\sigma(x_n^T w) - y_n) = X^T [\sigma(Xw) - y]$ (no closed form solution).
Hessian
 $H(w) = X^T S X$ with $S_{nn} = \sigma(x_n^T w) [1 - \sigma(x_n^T w)]$

8.3 Exponential family
General form
 $p(y|\eta) = h(y) \exp[\eta^T \psi(y) - A(\eta)]$ where
Cumulant
 $A(\eta) = \ln[\int_y h(y) \exp[\eta^T \psi(y)] dy]$
 $\nabla A(\eta) = \mathbb{E}[\psi(y)]$
 $\nabla^2 A(\eta) = \mathbb{E}[\psi \psi^T] - \mathbb{E}[\psi] \mathbb{E}[\psi^T]$
Link function
 $\eta = g^{-1}(\mu) \Leftrightarrow \mu = g(\eta)$

- $\eta_{\text{gaussian}} = (\mu/\sigma^2, -1/2\sigma^2)$
- $\eta_{\text{poisson}} = \ln(\mu)$
- $\eta_{\text{bernoulli}} = \ln(\mu/1 - \mu)$
- $\eta_{\text{general}} = g^{-1}(\frac{1}{N} \sum_{n=1}^N \psi(y_n))$

8.4 Nearest Neighbor Models
Performs best in low dimensions.

8.4.1 k-NN
 $f_{St,k}(x) = \frac{1}{k} \sum_{n: x_n \in \text{ngb}_{St,k}(x)} y_n$ Pick odd k so there is a clear winner. Large $k \rightarrow$ large bias small variance (inv.)

8.4.2 Error bound
 $\mathbb{E}[\mathcal{L}_{St}] \leq 2\mathcal{L}_{f^*} + 4c\sqrt{d}N^{-1/d+1}$

8.5 Support Vector Machines (SVM)
Logistic regression with hinge loss : $\min_w \sum_{n=1}^N [1 - y_n x_n^T w]_+ + \frac{\lambda}{2} \|w\|^2$ where $y \in [-1; 1]$ is the label and $\text{hinge}(x) = \max\{0, x\}$. Convex but not differentiable so need subgradient.
We can also use duality : $\mathcal{L}(w) = \max_{\alpha} G(w, \alpha)$. For SVM $\min_w \max_{\alpha \in [0, 1]^N} \sum \alpha_n (1 - y_n x_n^T w) + \frac{\lambda}{2} \|w\|^2$ differentiable and convex.
Can switch \max and \min when convex in w and concave in α . This can make the formulation simpler:
 $w(\alpha) = \frac{1}{\lambda} \sum \alpha_n y_n x_n = \frac{1}{\lambda} X^T \text{diag}(y) \alpha$ which yields the optimisation problem: $\max_{\alpha \in [0, 1]^N} \alpha^T \mathbf{1} - \frac{1}{2\lambda} \alpha^T Y X X^T Y \alpha$ The solution is sparse (α_n is the slope of the lines that are lower bounds to the hinge loss).

8.6 Kernel Ridge Regression
From duality $w^* := X^T \alpha^*$ where $\alpha^* := (K + \lambda I_N)^{-1} y$ and $K = XX^T = \phi^T(x) \phi(x) = \kappa(x, x')$ (needs to be PSD and symmetric).

9 Unsupervised Learning
9.1 K-means clustering
 $\min \mathcal{L}(z, \mu) = \sum_n \sum_k^K z_{nk} \|x_n - \mu_k\|_2^2$ with $z_{nk} \in \{0, 1\}$ (unique assignments: $\sum_k z_{nk} = 1$).
Algorithm (Coordinate Descent)

- $\forall n$ compute $z_n = \begin{cases} 1 & \text{if } k = \text{argmin}_j \|x_n - \mu_j\|^2 \\ 0 & \text{otherwise} \end{cases}$
- $\forall k$ compute $\mu_k = \frac{\sum_n z_{nk} x_n}{\sum_n z_{nk}}$

- Issues
- Heavy computation
 - Spherical clusters
 - Hard clusters

Probabilistic model
 $p(X|\mu, z) = \prod_n \mathcal{N}(x_n | \mu_k, I) = \prod_n \prod_k^K \mathcal{N}(x_n | \mu_k, I)^{z_{nk}}$

9.2 Gaussian Mixture Models
 $p(X|\mu, z) = \prod_n p(x_n | z_n, \mu_k, \Sigma_k) p(z_n | \pi) = \prod_n \prod_k^K [\mathcal{N}(x_n | \mu_k, \Sigma_k)]^{z_{nk}} \prod_k^K [\pi_k]^{z_{nk}}$ where $p_i k = p(z_n = k)$
Marginal likelihood: z_n are latent variables so they can be factored

out from the likelihood $p(x_n | \theta) = \sum \pi_k \mathcal{N}(x_n | \mu_k, \Sigma_k)$. (number of parameters reduced from $O(N)$ to $O(D^2 K)$).

9.3 EM
9.3.1 GMM
Initiaize $\mu^{(1)}, \Sigma^{(1)}, \pi^{(1)}$.

- E-step: Compute the assignments. $q_{kn}^{(t)} := \frac{\pi_k^{(t)} \mathcal{N}(x_n | \mu_k^{(t)}, \Sigma_k^{(t)})}{\sum_k^K \pi_k^{(t)} \mathcal{N}(x_n | \mu_k^{(t)}, \Sigma_k^{(t)})}$
- Compute Marginal Likelihood
- M-step: Update

$$\mu^{(t+1)} = \frac{\sum_n q_{kn}^{(t)} x_n}{\sum_n q_{kn}^{(t)}}$$

$$\Sigma^{(t+1)} = \frac{\sum_n q_{kn}^{(t)} (x_n - \mu^{(t+1)}) (x_n - \mu^{(t+1)})^T}{\sum_n q_{kn}^{(t)}}$$

$$\pi^{(t+1)} = \frac{1}{N} \sum_n q_{kn}^{(t)}$$

9.3.2 General
 $\theta^{(t+1)} := \text{argmax}_{\theta} \sum_n \mathbb{E}_{p(z_n | x_n, \theta^{(t)})} [\log p(x_n | \theta)]$
10 Matrix Factorizations
10.1 Prediction
Find $X \approx WZ^T$ where $W \in \mathbb{R}^{D \times K}$ and $Z \in \mathbb{R}^{N \times K}$ with $K \ll D, N$. Large $K \rightarrow$ overfitting. If $K \geq \max\{D, N\}$ trivial solution ($W = \mathbf{1}_D$ or $Z = \mathbf{1}_N$).
Quality of reconstruction (not jointly convex nor identifiable):
 $\mathcal{L}(W, Z) := \frac{1}{2} \sum_{(d,n) \in \Omega} [x_{dn} - (WZ^T)_{dn}]^2 =$

$$\sum_{(d,n) \in \Omega} f_{dn}(w, z)$$

Regularizer: $\Omega(W, Z) = \frac{\lambda_w}{2} \|W\|_{Frob}^2 + \frac{\lambda_z}{2} \|Z\|_{Frob}^2$

Optimisation with SGD (compute ∇_w for a fixed user d' and ∇_z for a fixed item n'). ALS (assume no missing ratings): $Z_*^T = (W^T W + \lambda_z I_K)^{-1} W^T X$
 $W_*^T = (Z^T Z + \lambda_w I_K)^{-1} Z^T X$

10.2 Text Representation
Factorize the co-occurence matrix to get each row forming a representation of a word (W) or a context word (Z) respectively.

10.2.1 GloVe

$$f_{dn} := \min\{1, (n_{dn}/n_{max})^\alpha\}, \alpha \in [0; 1]$$

10.2.2 Skipgram/CBOW

Binary classification to separate real word pairs from fake ones.

10.3 FastText

Supervised sentence-level BoW.

11 Dimensionality reduction

11.1 SVD

$\mathbf{X} = \mathbf{USV}^T$, with $\mathbf{X} : D \times N$, $\mathbf{U} : D \times D$ orthonormal, $\mathbf{V} : N \times N$ orthonormal, $\mathbf{S} : D \times N$ diagonal PSD, values in descending order ($s_1 \geq s_2 \geq \dots \geq s_D \geq 0$).

Reconstruction $\|\mathbf{X} - \hat{\mathbf{X}}\|_F^2 \geq \|\mathbf{X} - \mathbf{U}_K \mathbf{U}_K^T \mathbf{X}\|_F^2 = \sum_{i \geq K+1} s_i^2$ \forall rank- K matrix $\hat{\mathbf{X}}$ (i.e. we should compress the data by projecting it onto these left singular vectors.)

Truncated SVD: $\mathbf{U}_K \mathbf{U}_K^T \mathbf{X} = \mathbf{US}_K \mathbf{V}^T$

Application to MF: $\mathbf{U} = \mathbf{W}$ and $\mathbf{SV}^T = \mathbf{Z}^T$. Reconstruction limited by the rank- K of \mathbf{W}, \mathbf{Z} .

11.2 PCA

Decorrelate the data. Empirical mean before: $\mathbf{NK} = \mathbf{XX}^T = \mathbf{US}_D^2 \mathbf{U}^T$. After $\tilde{\mathbf{X}} = \mathbf{U}^T \mathbf{X}$: $\mathbf{N}\tilde{\mathbf{K}} = \tilde{\mathbf{X}}\tilde{\mathbf{X}}^T = \mathbf{S}_D^2$ (the components are uncorrelated).

Pitfalls: not invariant under scalings.

12 Neural Networks

The output at the node j in layer l is $x_j^{(l)} = \phi(\sum_i w_{i,j}^{(l)} x_i^{(l-1)} + b_j^{(l)})$

12.1 Representation power

Error bound $\leq \frac{(2Cr)^2}{n}$ where C is the smoothness bound, n the number of nodes. We can approximate any sufficiently smooth 2-dimensional function on a bounded domain (on average) with σ activation, "pointwise" with ReLU).

12.2 Learning

Problem is not convex but SGD is stable. Backpropagation: Let $\mathcal{L}_n = (y_n - f^{(L+1)} \circ \dots \circ f^{(1)}(\mathbf{x}_n^{(0)}))^2$.

Forward pass

$\mathbf{x}^{(0)} = \mathbf{x}_n$. For $l = 1, \dots, L+1$ $\mathbf{z}^{(l)} = (\mathbf{W}^{(l)})^T \mathbf{x}^{(l-1)} + \mathbf{b}^{(l)}$, $\mathbf{x}^{(l)} = \phi(\mathbf{z}^{(l)})$

Backward pass

$\delta^{(L+1)} = -2(y_n - \mathbf{x}^{(L+1)})\phi'(\mathbf{z}^{(L+1)})$ and $\forall l : \delta^{(l)} = (\mathbf{W}^{(l+1)})^T \delta^{(l+1)} \circ \phi'(\mathbf{z}^{(l)})$

Final pass

$$\frac{\partial \mathcal{L}_n}{\partial w_{i,j}^{(l)}} = \delta_j^{(l)} \mathbf{x}_i^{(l-1)}, \quad \frac{\partial \mathcal{L}_n}{\partial b_j^{(l)}} = \delta_j^{(l)}$$

12.3 Activations

- sigmoid $\phi(x) = 1 - \sigma(x)$

$$\tanh \frac{e^x + e^{-x}}{e^x + e^{-x}} = 2\phi(2x) - 1$$

- ReLU, Leaky ReLU ($\max\{ax, x\}$)

12.4 Convolutional Neural Nets

Convolution with filter f : $x^{(1)}[n, m] = \sum_{k,l} f[k, l] x^{(0)}[n-k, m-l]$. Filter is local so no need for fully connected layers. We can use same filter at every position: *weight sharing*. Learning: run backprop by computing different weights, then sum the gradients of shared weights.

12.5 Overfitting


Adding regularisation is equivalent to weight decay (by $(1-\eta\lambda)$). Can also use dataset augmentation, dropout.

13 Graphical Models

13.1 Bayes Nets

$p(X_1, \dots, X_D) = p(X_1)p(X_2|X_1)\dots p(X_D|X_1, \dots, X_{D-1})$. One node is a random variable, directed edge from X_j to X_i if X_j appears in the conditioning $p(X_i|\dots, X_j, \dots)$. The graph must be *acyclic*.

Conditional independence: $p(X, Y) = p(X)p(Y)$ or given Z $p(X, Y|Z) = p(X|Z)p(Y|Z)$.



- $p(X_1, X_2, X_3) = p(X_3)p(X_1|X_3)p(X_2|X_3)$: X_1 and X_2 are independent given X_3
- $p(X_1, X_2, X_3) = p(X_1)p(X_3|X_1)p(X_2|X_3)$: X_1 and X_2 are independent given X_3
- $p(X_1, X_2, X_3) = p(X_1)p(X_2)p(X_3|X_1, X_2)$: X_1

and X_2 are not independent given X_3

$X \rightarrow Y$ path blocked by Z if it contains a variable such that either

- variable is in Z and it is head-to-tail or tail-to-tail
- node is head-to-head and neither this node nor any of its descendants are in Z .

X and Y are D-separated by Z iff every path $X \rightarrow Y$ is blocked by Z .

X is conditionally independent of Y D-separated by Z . Independence is symmetric.

The Markov blanket of a node X_i is the set of parents, children, and co-parents of the node X_i (other parents of its children).

14 Quick maff

Chain rule $h = f(g(w)) \rightarrow \partial h(w) = \partial f(g(w)) \nabla g(w)$

Gaussian $\mathcal{N}(y|\mu, \sigma^2) \frac{1}{\sqrt{2\pi\sigma^2}} \exp(-\frac{(y-\mu)^2}{\sigma^2})$

Multivariate Gaussian $\mathcal{N}(y|\mu, \Sigma) = \frac{1}{\sqrt{(2\pi)^D \det(\Sigma)}} \exp(-\frac{1}{2}(y-\mu)^T \Sigma^{-1}(y-\mu))$

Bayes rule $p(x|y) = \frac{p(y|x)p(x)}{p(y)}$

Logit $\sigma(x) = \frac{\partial \ln[1+e^x]}{\partial x}$

Naming Joint distribution $p(x, y) = p(x|y)p(y) = p(y|x)p(x)$ where

- $p(x|y) \rightarrow$ likelihood
- $p(y) \rightarrow$ prior
- $p(y|x) \rightarrow$ posterior
- $p(x) \rightarrow$ marginal likelihood

Marginal Likelihood

$$p(\mathbf{X}|\alpha) = \int_{\theta} p(\mathbf{X}|\theta) p(\theta|\alpha) d\theta$$

Posterior probability \propto Likelihood \times Prior

Maximising over a Gaussian is equivalent to minimising MSE:

$$\beta_{MAP}^* = \operatorname{argmax}_{\beta} p(y|X, \beta) p(\beta) \Leftrightarrow \beta^* = \operatorname{argmin}_{\beta} \mathcal{L}(\beta)$$

Identifiable model $\theta_1 = \theta_2 \rightarrow P_{\theta_1} = P_{\theta_2}$

14.1 Algebra

$$(PQ + I_N)^{-1}P = P(QP + I_M)^{-1}$$

$$\sum_n (y_n - \beta^T \mathbf{x}_n)^2 = (\mathbf{y} - \mathbf{X}\beta)^T (\mathbf{y} - \mathbf{X}\beta)$$

$$\sum_j \beta^2 = \beta^T \beta$$

Unitary / orthogonal: $\mathbf{UU}^T = \mathbf{U}^T \mathbf{U} = \mathbf{I}$ and $\mathbf{U}^T = \mathbf{U}^{-1}$. Rotation matrix (preserves length of vector).

15 Mock Exam Notes

15.1 Normal equation

Unique if convex.

$$\frac{1}{\sigma_k^2} X(X^T w_k - y_k) + w_k = 0 \Leftrightarrow$$

$$w_k^* = (\frac{1}{\sigma_k^2} X X^T + I_D)^{-1} \frac{1}{\sigma_k^2} X y_k$$

15.2 MAP solution

$$\mathcal{L}(w) = \sum_k \sum_n \frac{1}{2\sigma_k^2} (y_{nk} - x_n^T w_k)^2 +$$

$$\frac{1}{2} \sum_k \|w_k\|^2 \rightarrow \text{Likelihood } p(y|X, w) =$$

$$\prod_n \prod_k \mathcal{N}(y_{nk} | w_k^T x_n, \sigma_k^2) \text{ and prior } p(w) = \prod_k \mathcal{N}(w_k | 0, I_D)$$

15.3 Convexity

$\ln[\sum_k e^{t_k}]$ is convex. Linear sum of parameters is convex.

15.4 Deriving marginal distribution

$$p(y_n | x_n, r_n = k, \beta) = \mathcal{N}(y_n | \beta_k^T \tilde{x}_n, 1)$$

Assume r_n follows a multinomial $p(r_n = k | \pi) = \pi_k$. Derive the marginal

$$p(y_n | x_n, \beta, \pi) = \sum_k p(y_n, r_n = k | x_n, \beta, \pi) =$$

$$\sum_k^K p(y_n | r_n = k, x_n, \beta, \pi) \cdot \pi_k =$$

$$\sum_k^K \mathcal{N}(y_n | \beta_k^T \tilde{x}_n, \sigma^2) \cdot \pi_k$$

15.5 MF

$$\hat{r}_{um} = \langle \mathbf{v}_u, \mathbf{w}_m \rangle + b_u + b_m \quad \mathcal{L} =$$

$$\frac{1}{2} \sum_u m (\hat{r}_{um} - r_{um})^2 + \frac{\lambda}{2} [\sum_u (b_u^2 + \|\mathbf{v}_u\|^2) + \sum_m (b_m^2 + \|\mathbf{w}_m\|^2)]$$

The optimal value for b_u for a particular user u' : $\sum_u m (\hat{r}_{um} - r_{um}) + \lambda b_{u'} = 0$.

Problem jointly convex? Compute $H(\hat{r}(v, w)) = \begin{bmatrix} 2w^2 & 4vw - 2r \\ 4vw - 2r & 2v^2 \end{bmatrix}$

which is not PSD in general.

16 Multiple Choice Notes

16.1 True statements

- Regularisation term sometimes renders the min. problem into a strictly concave/convex problem.

- k-NN can be applied even if the data cannot be linearly separated.

$$\max\{0, x\} = \max_{\alpha \in [0, 1]} \alpha x$$

$$\min\{0, x\} = \min_{\alpha \in [0, 1]} \alpha x$$

$$g(x) = \min_y f(x, y) \Rightarrow g(x) \leq f(x, y)$$

$$\max_x g(x) \leq \max_x f(x, y)$$

$$\max_x \min_y f(x, y) \leq \min_y \max_x f(x, y)$$

$$\nabla_W (\mathbf{x}^T \mathbf{W} \mathbf{x}) = \mathbf{x} \mathbf{x}^T$$

$$\nabla_x (\mathbf{x}^T \mathbf{W} \mathbf{x}) = (\mathbf{W} + \mathbf{W}^T) \mathbf{x}$$

K-means: optimal cluster (resp. centers) init \rightarrow one step optimal representation points (resp. clusters).

Logistic loss is typically preferred over L_2 loss in classification tasks.

For optimizing a MF of a $D \times N$ matrix, for large D, N : per iteration, ALS has an increased computational cost over SGD and per iteration, SGD cost is independent of D, N .

The complexity of backprop for a nn with L layers and K nodes/layer is $O(K^2 L)$

CNN where the data is laid out in a one-dimensional fashion and the filter/kernel has M non-zero terms. Ignoring the bias terms, there are M parameters.

16.2 Convex functions

- $f(x) = x^\alpha, x \in \mathbb{R}^+, \forall \alpha \geq 1$ or ≤ 0
- $f(x) = -x^3, x \in [-1, 0]$
- $f(x) = e^{ax}, \forall x, a \in \mathbb{R}$
- $f(x) = \ln(1/x), x \in \mathbb{R}^+$
- $f(x) = g(h(x)), x \in \mathbb{R}, g, h$ convex and increasing over \mathbb{R}
- $f(x) = ax + b, x \in \mathbb{R}, \forall a, b \in \mathbb{R}$
- $f(x) = |x|^p, x \in \mathbb{R}, p \geq 1$
- $f(x) = x \log(x), x \in \mathbb{R}^+$

16.3 Non-convex functions

$$f(x) = x^3, x \in [-1, 1]$$

$$f(x) = e^{-x^2/2}, x \in \mathbb{R}$$

$$\sum \mathcal{N}$$

$$\sin(x) \forall x \in \mathbb{R}$$